



From Censorship to Detoxification: Rewriting Harmful Memes

Alex Dell'Orto, Christian Garzone, Dario Liuzzo

Group 31 EE-559: Deep Learning, EPFL, 2026



Problem definition

Harmful memes combine images and overlaid text to spread offensive content across both online platforms and physical spaces. Existing moderation relies mainly on **censorship** after exposure: removing, blurring, or suppressing the content.

This project instead proposes **text detoxification**: rewriting harmful meme text in real time while preserving the core meaning and visual coherence of the original meme. The goal is proactive moderation without erasure, enabling safer message delivery rather than simple deletion.

The research hypothesis is that a large multimodal teacher can explain why a meme is harmful and generate safer pseudo-rewrites, then transfer this context-aware rewriting behavior to a compact BART-large student for efficient deployment.

Related works

- **Hateful meme detection.** The dominant literature frames harmful memes as a classification task. The Hateful Memes challenge showed that multimodal fusion models outperform unimodal baselines, and later VisualBERT- and CLIP-based systems further improved image–text reasoning [1, 2, 3]. These methods detect or remove content, but do not preserve the original communicative intent through rewriting.
- **Text detoxification and style transfer.** ParaDetox treats toxicity as a style attribute and trains sequence-to-sequence models such as BART/T5 to neutralize toxic sentences [4]. However, plain-text detoxification ignores the image context that often gives meme text its harmful meaning.
- **Our gap.** We use LLaVA-Next’s multimodal understanding as a teacher to produce target groups, visual evidence, implicit meanings, and pseudo-detoxified rewrites [5]. These labels are distilled into a lightweight BART-large student, enabling context-aware rewriting at much lower inference cost.

Method

Stage 0 – Data Preparation
EasyOCR extraction + CLIP meme filtering

Stage 1 – LLaVA-Next Teacher
Hate explanations + pseudo-detoxified rewrites

Stage 2 – BART-large Student
LoRA fine-tuning for meme text detoxification

Stage 3 – Proxy Network
CLIP features → BART soft memory for cheaper inference

Evaluation
Text STA, Δ toxicity, SIM, CLIP, CO₂

- **Stage 0:** EasyOCR extracts meme text; CLIP filters non-meme-like images, using stricter negative prompts for MMHS150K.
- **Stage 1:** LLaVA-Next runs in two sharded passes: first structured hate explanations, then pseudo-detoxified rewrites.
- **Stage 2:** noisy pseudo-labels are filtered by parse validity, Stage 1 quality checks, positive toxicity reduction, and text quality before LoRA fine-tuning BART-large.
- **Stage 3:** a proxy MLP maps CLIP image/text features to a short BART encoder-memory sequence, exploring a VLM-free deployment path.

Student LoRA setup

Conditions

facebook/bart-large, 406M parameters.
Rank 32, alpha 64, dropout 0.05; 17M trainable parameters (~4.3%).
full, target_only, visual_only, and none.

Datasets	
HarMeme	~3,500 memes related to COVID-19 and US politics, manually labelled as harmful/non-harmful with associated textual claims [6].
MAMI	10,000 misogynist memes from Twitter [7].
MMHS150K	~150,000 heterogeneous and noisy Twitter image–text posts [8].

All datasets are filtered to retain meme-like images only. The final Stage 2 supervision contains 3,578 clean pseudo-rewrite pairs from 7,918 loaded LLaVA rows, split into 3,220 train and 358 validation examples. Combining curated and noisy sources exposes the model to diverse toxicity types and improves generalization pressure.

Validation

Quality on 358 validation memes. Text STA is the mean non-toxic probability from a RoBERTa toxicity classifier; Δ toxicity is the improvement over the original meme text; SIM is BERTScore F1; CLIP measures image–text alignment.

System	Text STA	Δ toxicity	SIM	CLIP
LLaVA-Next Teacher	0.9517	0.2804	0.3938	0.6253
DetoxLLM baseline	0.8894	0.2182	0.4202	0.6251
BART-large Full	0.9261	0.2548	0.2991	0.6236
Proxy + BART	0.8084	0.1372	0.5527	0.6312

Reading the trade-off. BART-large Full gives the strongest student toxicity reduction and improves over DetoxLLM [9] on Text STA and Δ toxicity. DetoxLLM and Proxy + BART preserve more surface similarity, but the proxy currently sacrifices detoxification strength. This matches the project goal: distil a useful portion of LLaVA’s multimodal rewriting behavior into a much smaller student.

Compute and emissions.

Pipeline component	GPU-time	Energy	CO ₂
Stage 0 – OCR + CLIP filtering	2.8 h	0.76 kWh	26.7 g
Stage 1A – LLaVA explanations	14.1 h	6.61 kWh	230.3 g
Stage 1B – LLaVA rewrites	14.5 h	5.78 kWh	201.3 g
Stage 2 – BART LoRA training	2.1 h	0.82 kWh	28.4 g
BART finetuned inference	12.3 min	0.04 kWh	1.5 g
Proxy inference	3.4 min	0.01 kWh	0.5 g

Stage 1 accounts for most emissions because the 7B teacher is used to create pseudo-labels. This is a one-time distillation cost: the deployed BART student is 76.6× faster than the LLaVA teacher and 16.3× faster than DetoxLLM in the repository benchmark.

+3.67 pts
Text STA vs. DetoxLLM

76.6×
BART speedup vs. LLaVA

496 g
total measured CO₂

Limitations

- Dataset noise and heterogeneity remain substantial, especially in MMHS150K.
- The student depends on LLaVA pseudo-labels; filtering improves precision but reduces the training set size.
- BART has limited direct visual grounding at inference time; image context is distilled through text fields rather than consumed as pixels.
- Evaluation is only partially multimodal. VisualBERT scores were difficult to interpret as absolute detoxification metrics in this run.
- The proxy model still shows an unbalanced toxicity–similarity trade-off.

Conclusion

This project shows that multimodal supervision from LLaVA-Next can be distilled into a lightweight BART-large student that performs meme-specific text detoxification: it reduces toxicity while preserving semantic alignment and image–text coherence.

The result supports a shift from reactive censorship to constructive rewriting. Future work should improve dataset quality, strengthen direct multimodal grounding, and integrate OCR plus inpainting modules for live rewriting of harmful text in video streams and real-world physical environments.

References

[1] Douwe Kiela, Hamed Firooz, Aravind Mohan, Vedanuj Goswami, Amanpreet Singh, Pratik Ringshia, and Davide Testuggine. The hateful memes challenge: Detecting hate speech in multimodal memes. In *Advances in Neural Information Processing Systems*, 2020.

[2] Liunan Harold Li, Mark Yatskar, Da Yin, Cho-Jui Hsieh, and Kai-Wei Chang. Visualbert: A simple and performant baseline for vision and language. In *arXiv preprint arXiv:1908.03557*, 2019.

[3] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, 2021.

[4] Varvara Logacheva, Daryna Dementieva, Sergey Ustyantsev, Daniil Moskovskiy, David Dale, Irina Krotova, Nikita Semenov, and Alexander Panchenko. Paradox: Detoxification with parallel data. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, 2022.

[5] Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Llava-next: Improved reasoning, ocr, and world knowledge. *arXiv preprint*, 2024.

[6] Shraman Pramanick, Shivam Sharma, Dimitar Dimitrov, Md. Shad Akhtar, Preslav Nakov, and Tanmoy Chakraborty. Harmeme: A dataset for hate speech detection in memes. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, 2021.

[7] Elisabetta Fersini, Francesca Gasparini, Giulia Rizzi, Aurora Saibene, Berta Chulvi, Paolo Rosso, Alyssa Lees, and Jeffrey Sorensen. Semeval-2022 task 5: Multimedia automatic misogyny identification. In *Proceedings of the 16th International Workshop on Semantic Evaluation*, 2022.

[8] Raul Gomez, Jaume Gibert, Lluís Gomez, and Dimosthenis Karatzas. Exploring hate speech detection in multimodal publications. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision Workshops*, 2020.

[9] Skyler Hallinan, Ameet Deshpande, Sourmyadeep Sarkar, Nitesh Goyal, Kai-Wei Chang Yang, and Daniel Khoshabi. Detoxllm: An llm-based text detoxification system. UBC-NLP system report, 2023.