




Inference path: CLIP extracts image/text features, the MLP predicts K BART-sized soft memory tokens, and these tokens are prepended to the text encoder states before BART decoding.